

Cluster-Computing and Parallelisation for the Multi-Dimensional PH-Index

Master Thesis

Bogdan Aurel Vancea

<bvancea@student.ethz.ch>

Prof. Dr. Moira C. Norrie
Tilmann Zaeschke
Christoph Zimmerli

Global Information Systems Group
Institute of Information Systems
Department of Computer Science
ETH Zurich

1st December 2014



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich



Abstract

Here comes the abstract.

Contents

1	Introduction	1
1.1	Multi-dimensional Indexes	1
1.1.1	Background information	1
2	Design and Implementation	3
2.1	Algorithms	3
2.1.1	Requirements	3
2.1.2	Data partitioning	3
2.1.3	Data balancing	4
2.1.4	Basic Index operation	4
2.1.5	Iterators	4
2.1.6	Range search	4
2.1.7	K Nearest neighbours	5
2.2	Implementation	5
3	Performance Analysis	7
3.1	Benchmark	7
4	Conclusions	9

1

Introduction

1.1 Multi-dimensional Indexes

1.1.1 Background information

Need to add an introduction here.

This is an example of how to cite a scientific publication [1] from your bibliography (BibTeX¹ file). And this example shows how you create links within your documents, e.g. link to section 1.1.

¹<http://en.wikipedia.org/wiki/BibTeX>

2

Design and Implementation

2.1 Algorithms

2.1.1 Requirements

The distributed index has the following requirements:

1. Eventual consistency. The index is eventual consistent, meaning that the following assumptions are true:
 - (a) Any entries that are inserted at some point t in time will eventually be returned in results at a following point in time $t + \epsilon$.
 - (b) In the case that a client starts an insert operation at time t and another client, in a separate thread or a separate, issues a query at time $t + \epsilon$, the response of the query of the second client is not guaranteed to contain the point inserted by the first client.
2. Re-balancing transparency. The balancing process should be transparent with respect to the client. This means that the client API should not know that there might be multiple versions of the key mapping available due to any on-going balances.

2.1.2 Data partitioning

Describe here how the key-value pairs are partitioned across the hosts.

2.1.3 Data balancing

Describe how the key-value pairs are balanced across the index nodes. The cluster should be able to properly balance the amount of keys that are stored in each server.

A simple load balancing strategy would be the following:

1. Upon reaching a certain threshold t of keys stored, a host decides it has to split its zone.
2. This host first sends a broadcast to all nodes to request the number of keys they all store. It then chooses the hosts with the fewest keys and considers that this key is the split receiver. The splitter decides to split its zone in half (or in such manner that around half of the keys held in the initial zone are moved to the receiver).
3. The splitting host sends the keys to the receiver. (What does the receiver do with them? Duplicates could appear in case of KNN or range queries.) .
4. After the receiver received all keys, the mapping is updated on the Zookeeper.
5. The splitter then deletes all of the keys sent from its own index (maybe this can be done really fast by removing a bunch of nodes from the index).

Ideas that might not work:

- Storing information like the number of keys held by each node in Zookeeper. This would mean that the ZooKeeper would need to be notified on each insert, which would cause severe scalability and availability issues. What could work is that the ZK is notified when each host reaches a certain key threshold (i.e, every 1000 keys inserted, a request is sent to ZK).

2.1.4 Basic Index operation

Describe the point operations here. These are operations that affect a single key-value pair, like get, put, delete, contains, etc.

2.1.5 Iterators

Describe how iterators are handled when dealing with a cluster of index servers. Describe the current algorithm used and the alternatives.

2.1.6 Range search

Describe how the range search is performed. Describe how the number of hosts that need to be queried is reduced and what the alternatives are.

2.1.7 K Nearest neighbours

The K nearest neighbours should be found in the following manner:

1. First, a request for the k nearest neighbours is sent to the host that holds the query key. This query will return a number of candidate points.
2. After finding the furthest neighbour fn from the set of candidates obtained at the previous step one could do the following:
 - look into all neighbouring areas and check for neighbours that are closer to q than fn .
 - find all zones intersecting the square $(q - \text{dist}(q, fn), q + \text{dist}(q, fn))$ and perform a knn query into those areas. Then apply an additional knn to combine candidates.
 - find all zones intersecting the square $(q - \text{dist}(q, fn), q + \text{dist}(q, fn))$ and perform a range query in those areas. Then apply an additional knn to combine candidates.
 - find all zones intersecting the square $(q - \text{dist}(q, fn), q + \text{dist}(q, fn))$ and perform a range query followed by a filtering based on $\text{dist}(q, fn)$ in those areas. Then apply an additional knn to combine candidates.

This algorithm works under that assumption that all hosts hold at least k points. If that is not the case, and fewer candidates are returned from the first query, one would need to increase the query space and check all of the neighboring areas of the zone holding the query key. To achieve correctness even if all the cluster holds fewer than k keys, one could iteratively query more neighbours in each steps (i.e. first query the 1 - hop neighbourhood, then the 2 - hop neighbourhood) and stop after all of the hosts have been queried.

2.2 Implementation

Describe the technologies used, the reasons for which these technologies were chosen and any alternatives.

Currently used frameworks:

ZooKeeper ZooKeeper is used for stored cluster metadata and membership. Currently, the only alternative would have been to implement such a distributed storage manager manually. Using ZooKeeper saved a lot of development time.

Netty Netty is a Java IO library and it is used to implement the server request handling component. Alternatives would have been Java NIO library.

Kryo Kryo is very fast serialization library for Java and it is used to serialize the values that have to be stored on the server. These objects need to be transformed into a representation that can be sent over the network. Kryo is faster than the Java serialization, does not require the implementation of the Serializable interface and transforms the objects into byte arrays. This should make the representation smaller than simply transforming the object to a string.

3

Performance Analysis

3.1 Benchmark

4

Conclussions

List of Figures

List of Tables

Acknowledgements

Bibliography

- [1] Alfonso Murolo. Designing wordpress themes by example. Master's thesis, ETH Zurich, 2013.